



## CPE 412/512

### Fall Semester 2017

#### Using the Batch Queuing System on the dmc.asc.edu system using *run\_script* utility for MPI programs

*The Batch Queuing system is the method by which students should execute programs on the Alabama Supercomputer DMC system. It allows your program to have exclusive access to the specified portion of the system.*

*To invoke the queuing system for MPI programs follow the following steps.*

**Step 1:** Change to your appropriate working subdirectory under your home directory

**Step 2:** Create the source file and compile for mpi in the usual manner.

ex. **mpic++ hello\_world\_MPI.cpp -o hello\_world\_MPI**

This may require that you load the mpi environment using the

**module load openmpi**

command before executing the **mpic++** command.

**Step 3:** Then using your favorite text editor (vi, nano, emacs, pico, etc.) create a bash script file that contains the **module load** command for openmpi and at least one **srun** command to start the MPI system. The **srun** command functions in a similar manner to the **mpirun** and **mpiexec** commands but does not require a separate parameter to be entered as part of the command to specify the number of MPI processes that will be generated. With the **srun** command the number of MPI processes generated is automatically set to match the number of processing cores that was specified by the user when the job was submitted to the queuing system. For example if you want to have a two process MPI run for a file with an executable name "*hello\_world\_MPI*" you may create the file "*hello\_world\_MPI.sh*" which contains the following lines of text.

```
#!/bin/bash
```

```
module load openmpi
```

```
srun hello_world_MPI
```

In this example, the number of MPI processes can then be set to 2 by requesting that two processing cores be used when this MPI program is executed by the batch queuing system. This is done in Step 5.

**Step 4:** Make sure that your execution script file has execute privileges

```
chmod 744 hello_world_MPI.sh
```

**Step 5:** To invoke the batch queuing system use the *run\_script* command as the command line argument the name of your file. In this case you would type **run\_script hello\_world\_MPI.sh**.

The **run\_script** command is interactive requiring you to respond to a number of questions as shown below:.

```

Queue           Wall Time      Mem  # Cores
-----
express         4:00:00      16gb   1-4
small           60:00:00       4gb   1-8
medium          150:00:00      16gb   1-16
large           360:00:00     120gb   1-64
huge            360:00:00     256gb  64-128
class           16:00:00      64gb   1-64

```

Enter Queue Name (default <cr>: small) **class** ← enter class to select "class" queue

Enter number of processor cores (default <cr>: 1 ) **2** ← enter 2 to reserve two processing cores

Enter Time Limit (default <cr>: 16:00:00 HH:MM:SS) ← accept defaults press <enter>

Enter memory limit (default <cr>: 2gb ) ← accept defaults press <enter>

Enter a name for your job (default: helloworldMPIshSCR)

Run on; uv, dmc, nehalem, ivy-bridge, etc (default: any)

**dmc** ← enter dmc

```

=====
==== Summary of your script      job      =====
=====
The script file is: hello_world_MPI.sh
The time limit is 16:00:00 HH:MM:SS.
The memory limit is: 2gb
The job will start running after: 2016-09-14T20:35:53
Job Name: helloworldMPIshSCR
Virtual queue: class
QOS: -p dmc,uv,knl --qos=class
Constraints: --constraint=dmc
Queue submit command:
sbatch -p dmc,uv,knl --qos=class -J helloworldMPIshSCR --begin=2016-09-14T20:35:53 --
requeue --mail-user=wellsbe@uah.edu -o helloworldMPIshSCR.o%A -t 16:00:00 -N 1-1 -n 2
--mem-per-cpu=1000mb --constraint=dmc
Batch Queue
Job ID
Submitted batch job 49509

```

For this example you would answer the questions in the following manner. Enter **class** for *Queue Name*, **2** (in this case) for number of *processor cores* (will be used by the *srunk* command(s) in the script file), and **dmc** for the machine it is to run on. Press the <Enter> key to accept all other defaults for the other prompts. All screen output should go to the file *helloworldMPIshSCR.o49509* (i.e. the <filename>SCR.o<job number> file in general).

You can view your status using the *squeue*, *squeue1*, *squeue2*, *squeue3*, or *squeue4* commands which vary from one another in terms of the information format and job/queue attributes that are shown. For example, you can type

### squeue4

to view the status of your job(s). If you are user *uahcls01* then the following output could be generated.

JobID	User	QOS	JobName	ReqCPUS	ReqMem	Timelimit	State	Elapsed	Partition
49509	uahcls01	class	helloworl+	2	1000Mc	16:00:00	RUNNING	00:00:13	dmc-ivy-b+

Unlike program jobs run without this queue management in place, jobs run using the batch queue will be the only ones executed on the portion of the system you have reserved while your program is running. If your job run time exceeds the *Timelimit* then your job will be terminated and dequeued from the batch system. The *Timelimit* (hours:minutes:seconds) represents the maximum time product of parallel execution time multiplied by the number of processing cores that you specified when the job entered the queuing system. In this example the default value was taken when the queue was invoked.

You can also remove a job from the queue before it completes execution by using the *scancel* command with your job number as the command line argument. For example if you realize you made a mistake in your code and do not want job number 49509 run after it has been submitted to the queue, simply type

**scancel 49509**

to delete it.