



# Parallel Brute Force Dictionary Attack Using OpenMPI and OpenMP

CPE 512

KYLE RAY

December 8, 2017

## Contents

Abstract .....	2
Introduction .....	3
Environment .....	3
Test Layout .....	4
Results .....	5
MPI Results.....	6
OpenMP Results.....	8
Conclusion.....	10
Future Work.....	10
Appendix .....	11
MPI Source Code.....	11
OpenMP Source Code.....	17

## Abstract

In this experiment, a popular brute force attack known as the Dictionary Attack was implemented and then parallelized. The parallelization was performed using a message-passing interface (MPI) to distribute the dictionary amongst a cluster of computers and an API known as OpenMP which utilizes compiler pragmas as well as interface calls to parallelize existing serial code. Also, three different hashing algorithms (MD5, SHA1, SHA256) were tested in this experiment to see the timing affects a more complex algorithm might have on finding the password. It was observed that the MPI implementation performed poorly for each hashing algorithm in most cases worse than the serial implementation; whereas, the OpenMP implementation had great speed up and efficiency. The experiment shows that while parallelization does make an impact on performing this attack, the method that is employed is very important.

## Introduction

In this day and age most of our daily life is performed in the digital realm. To protect our digital identity, which can grant access to everything from our bank accounts to our social security number, we must have a secure method to access this information. Passwords have been used since the earliest days of computing and there have been many techniques developed to keep these passwords secure. In this project, I would like to explore a form of securing passwords known as hashing and one of the attacks used to crack it.

Hashing is a one-way mathematical transformation of an input to a constant sized output string. Hence a user gives a password, which then passes through a “hashing” algorithm and is eventually stored. There are many different hashing algorithms out there which range in complexity. Some of the most common hashing techniques are MD5, SHA-1, SHA-256, and so on.

Example MD5 Hashed Password:

Password = “Parallel\_Programming\_is\_fUn”

Hash = 6243ebf30a0642f7dadoec17cc916df9

The purpose of this experiment was to determine the benefits that a hacker would have if they were to implement a parallel version of a very well-known penetration attack. The penetration attack used is the Brute Force Dictionary Attack. The attack assumes that the hacker possesses some password hashes from the victim and is now trying to determine the human readable password via confirming the password that created the hash. To do this the hacker will use a known Dictionary of either previously cracked passwords or just popular words in general, apply the hashing algorithm to each word and check each hash against the victim’s hash. Once a matching hash is found the hacker will more than likely have access to any of the victim’s accounts using this password.

Performing this attack can be quite a process intensive task if the dictionary is large and parallelism can be utilized to break these tasks down into smaller subtasks and thus decrease the overall time it takes to successfully find the password.

## Environment

The experiment was carried out on the Jetson cluster located at the University of Alabama in Huntsville.

## Test Layout

To perform the experiment a popular dictionary of passwords known as rockyou.txt was incorporated. The list was trimmed down from around fifteen million passwords to just ten million, removing most of the passwords with spaces. The dictionary can be found on the [SkullSecurity](#) website.

The three hashing algorithms ([MD5](#), [SHA1](#), [SHA256](#)) were taken from zedwood.com, license file included with source code. See appendix for implementation source code.

The parallel implementations were both run with the following conditions. They programs were tasked each with finding the words telva, khridaniel, and arisfarias which are near the front, middle, and end respectively of the dictionary. The intent of the test is to simulate three possible searches starting from the beginning of the list. This task was run on each implementation with 1, 2, 4, and 8 processes/threads. The times were recorded and checked against a base serial version to determine the overall speed up and efficiency of each implementation.

The application begins by reading in the dictionary file line by line and storing it in a container. In the MPI case the program will scatter the list as evenly as possible to the number of processes desired and each process will perform the attack using the hashing algorithm specified. The OpenMP version utilizes the work-sharing construct for the for loop to divide the work amongst the number of threads specified. Once the password is found it is displayed to the user along with the hashed representation. The time to find the password is also displayed.

Below is an example output from each of the implementations.

### MPI Implementation:

```
kr0010@jetson:~/jetson/Project/MPI_Run$ mpirun -np 4 OpenMPI_Password_Crack 0
rockyou.txt telva
Reading Dictionary
Finished reading Dictionary with 10000000 passwords!
Starting the timer.
Scattering...
Pass found by process 1
Password is telva
Hash is bb95c8f4455a2385e3c5af1a301c6240
CLEAN UP: Process 1 is finished processing
CLEAN UP: Process 2 is finished processing
CLEAN UP: Process 3 is finished processing
CLEAN UP: Process 0 is finished processing
Time to find the password is 18.0492 seconds
Total execution time: 50.8507 seconds
```

### OpenMP Implementation:

```
kr0010@jetson:~/jetson/Project/OpenMP_Run$ ./OpenMP_Password_Crack 0
rockyou.txt telva 4
Reading Dictionary
```

```

Finished reading Dictionary with 10000000 passwords!
Starting the Timer
Processing...
Pass found by process 3
Password is telva
Hash is bb95c8f4455a2385e3c5af1a301c6240
Time to find the password is 4.48877 seconds

```

The MPI version was originally implemented in a way that would stop processing after the password had been found by a process, this was achieved by keeping a flag and having every process check that flag to see if they needed to stop. In the end this implementation worked but was dreadfully slow. The blocking communication and the time to communicate are far too costly to be utilized in this sense especially because the loops are performing at least  $listSize/numtasks$  with the list size being ten million. The final decision was to remove this and just keep the execution time of the process that found the password because in the end, if a hacker were to use this method, once the password is output to the console the job is done.

## Results

This portion of the report contains the data collected when running the experiment. The table below shows the serial run time of each test case.

**Table 1: Serial Run Times for Each Test Scenario**

Serial MD5	Run Time (s)
Front of Dictionary (telva)	19.7664
Middle of Dictionary (khrisdaniel)	34.8885
End of Dictionary (arisfarias)	56.9507
Serial SHA1	Run Time (s)
Front of Dictionary (telva)	29.3506
Middle of Dictionary (khrisdaniel)	64.5736
End of Dictionary (arisfarias)	99.7225
Serial SHA256	Run Time (s)
Front of Dictionary (telva)	36.5515
Middle of Dictionary (khrisdaniel)	75.9972
End of Dictionary (arisfarias)	110.581

The algorithms are ordered by complexity with SHA256 being the most complex hashing algorithm tested.

## MPI RESULTS

This section houses the data collected from the MPI test runs.

**Table 2: MPI Run Times for Each Test Scenario**

MPI MD5	NP = 1	NP = 2	NP = 4	NP = 8
Front of Dictionary (telva)	23.0993	44.5034	24.8596	21.1322
Middle of Dictionary (khrisdaniel)	43.7216	36.4375	27.1765	19.9634
End of Dictionary (arisfarias)	63.7789	61.895	34.5734	16.8332
MPI SHA1	NP = 1	NP = 2	NP = 4	NP = 8
Front of Dictionary (telva)	40.2165	113.334	45.8712	31.0657
Middle of Dictionary (khrisdaniel)	77.8335	96.8681	64.3992	38.7767
End of Dictionary (arisfarias)	111.519	130.166	60.6241	31.4162
MPI SHA256	NP = 1	NP = 2	NP = 4	NP = 8
Front of Dictionary (telva)	43.5325	111.886	46.3524	41.2534
Middle of Dictionary (khrisdaniel)	80.7348	94.4664	64.7486	32.6261
End of Dictionary (arisfarias)	113.118	129.271	75.7362	41.5177

Comparing the results from Table 2 with the results from Table 1 we can see that the MPI version performed poorly in the cases where less than eight processes were used and that the word was found near the middle or end of the file. In all other cases the serial version performed better. The values in Table 2 **DO** take into account the scattering of the dictionary to the other processes, so this incurs quite a bit of communication overhead as well as the overhead of creating the processes. Refer to the next two tables for the speed up and efficiency of this implementation.

**Table 3: MPI Speed Up for Each Test Scenario**

MPI MD5 S(p)	NP = 1	NP = 2	NP = 4	NP = 8
Front of Dictionary (telva)	0.855714242	0.444154829	0.795121402	0.935368774
Middle of Dictionary (khrisdaniel)	0.797969425	0.957488851	1.283774585	1.74762315
End of Dictionary (arisfarias)	0.892939514	0.920117942	1.647240364	3.383236699
MPI SHA1 S(p)	NP = 1	NP = 2	NP = 4	NP = 8
Front of Dictionary (telva)	0.729814877	0.258974359	0.639848096	0.944791201
Middle of Dictionary (khrisdaniel)	0.829637624	0.666613674	1.002708108	1.66526806
End of Dictionary (arisfarias)	0.894219819	0.76611788	1.644931636	3.174238132
MPI SHA256 S(p)	NP = 1	NP = 2	NP = 4	NP = 8
Front of Dictionary (telva)	0.839637053	0.326685197	0.788556795	0.88602394
Middle of Dictionary (khrisdaniel)	0.941318985	0.804489215	1.173727308	2.329337555
End of Dictionary (arisfarias)	0.977572093	0.855420009	1.460081177	2.663466425

Although the MPI implementation reduced the execution in a few cases the speed up achieved was not desirable compared to the number of resources used and this can be seen by viewing the efficiency table below.

**Table 4: MPI Efficiency for Each Test Scenario**

MPI MD5 E(p)	NP = 1	NP = 2	NP = 4	NP = 8
Front of Dictionary (telva)	0.855714242	0.222077414	0.19878035	0.116921097
Middle of Dictionary (khrisdaniel)	0.797969425	0.478744425	0.320943646	0.218452894
End of Dictionary (arisfarias)	0.892939514	0.460058971	0.411810091	0.422904587
MPI SHA1 E(p)	NP = 1	NP = 2	NP = 4	NP = 8
Front of Dictionary (telva)	0.729814877	0.129487179	0.159962024	0.1180989
Middle of Dictionary (khrisdaniel)	0.829637624	0.333306837	0.250677027	0.208158508
End of Dictionary (arisfarias)	0.894219819	0.38305894	0.411232909	0.396779766
MPI SHA256 E(p)	NP = 1	NP = 2	NP = 4	NP = 8
Front of Dictionary (telva)	0.839637053	0.163342599	0.197139199	0.110752992
Middle of Dictionary (khrisdaniel)	0.941318985	0.402244608	0.293431827	0.291167194
End of Dictionary (arisfarias)	0.977572093	0.427710005	0.365020294	0.332933303

Looking at the table above we can see that in most cases where more than one process is used that the efficiency declines significantly.

It seems that, from this experiment, the process of hashing a string and doing string compares doesn't really suit the type of parallelism offered by using MPI.



## OPENMP RESULTS

Not only was the OpenMP version much easier to implement, the results achieved were vastly better than the MPI version implemented. This implementation achieved a reasonable decrease in execution time every time more threads were utilized.

**Table 5: OpenMP Run Times for Each Test Scenario**

OpenMP MD5	NT = 1	NT = 2	NT = 4	NT = 8
Front of Dictionary (telva)	17.6082	8.92389	6.66862	6.08179
Middle of Dictionary (khrisdaniel)	36.9943	18.1401	12.3859	10.8101
End of Dictionary (arisfarias)	57.7918	27.8238	13.8588	13.7383
OpenMP SHA1	NT = 1	NT = 2	NT = 4	NT = 8
Front of Dictionary (telva)	35.398	17.1591	9.55503	9.41605
Middle of Dictionary (khrisdaniel)	69.4872	35.6905	19.5332	19.4156
End of Dictionary (arisfarias)	106.239	51.7842	27.6193	27.9677
OpenMP SHA256	NT = 1	NT = 2	NT = 4	NT = 8
Front of Dictionary (telva)	37.0772	17.312	8.97647	8.78002
Middle of Dictionary (khrisdaniel)	76.7078	37.414	27.3433	20.387
End of Dictionary (arisfarias)	112.131	56.9685	39.4119	26.3841

Comparing Table 5 with the serial run times of Table 1 we can clearly see that this implementation is much faster. It's interesting, but makes sense, that the execution time for the case that the word is near the front and the application utilizes four and eight threads is near the same for each implementation. It seems that there is a limit for this case and adding more threads will not be as efficient.

**Table 6: OpenMP Speed Up for Each Test Scenario**

OpenMP MD5 S(p)	NT = 1	NT = 2	NT = 4	NT = 8
Front of Dictionary (telva)	1.122567895	2.214998168	2.964091521	3.250095778
Middle of Dictionary (khrisdaniel)	0.943077717	1.923280467	2.816791674	3.227398451
End of Dictionary (arisfarias)	0.985446032	2.046834005	4.109352902	4.145396446
OpenMP SHA1 S(p)	NT = 1	NT = 2	NT = 4	NT = 8
Front of Dictionary (telva)	0.82915984	1.710497637	3.071743364	3.117082004
Middle of Dictionary (khrisdaniel)	0.929287696	1.809265771	3.305838265	3.325861678
End of Dictionary (arisfarias)	0.938661885	1.925732173	3.610609248	3.565631067
OpenMP SHA256 S(p)	NT = 1	NT = 2	NT = 4	NT = 8
Front of Dictionary (telva)	0.985821475	2.111338956	4.071923596	4.163031519
Middle of Dictionary (khrisdaniel)	0.990736275	2.031250334	2.779371912	3.727728454
End of Dictionary (arisfarias)	0.986176882	1.941090252	2.805776935	4.191198487

As far as speed up, looking at the table above, the cases using two and four threads have close to ideal speed up compared to the other scenarios.

**Table 7: OpenMP Efficiency for Each Test Scenario**

OpenMP MD5 E(p)	NT = 1	NT = 2	NT = 4	NT = 8
Front of Dictionary (telva)	1.122567895	1.107499084	0.74102288	0.406261972
Middle of Dictionary (khrisdaniel)	0.943077717	0.961640234	0.704197919	0.403424806
End of Dictionary (arisfarias)	0.985446032	1.023417003	1.027338226	0.518174556
OpenMP SHA1 E(p)	NT = 1	NT = 2	NT = 4	NT = 8
Front of Dictionary (telva)	0.82915984	0.855248818	0.767935841	0.38963525
Middle of Dictionary (khrisdaniel)	0.929287696	0.904632886	0.826459566	0.41573271
End of Dictionary (arisfarias)	0.938661885	0.962866087	0.902652312	0.445703883
OpenMP SHA256 E(p)	NT = 1	NT = 2	NT = 4	NT = 8
Front of Dictionary (telva)	0.985821475	1.055669478	1.017980899	0.52037894
Middle of Dictionary (khrisdaniel)	0.990736275	1.015625167	0.694842978	0.465966057
End of Dictionary (arisfarias)	0.986176882	0.970545126	0.701444234	0.523899811

Overall OpenMP is the most efficient implementation as well as very easy to implement compared to the MPI version.

## Conclusion

After the experiment it is clear to see that a hacker could benefit from using parallel processing to perform the brute force dictionary attack. The experiment also demonstrated that some tools are better for the job than others in the case where MPI was very tedious to use because of all the synchronization and message passing and OpenMP involved changing a few lines of code in the serial version to achieve a substantial performance increase.

## Future Work

Optimization is a never-ending endeavor. The parallel implementations in this experiment were nowhere near totally optimized and there is a lot of room for improvement in this area. Also, the actual searching algorithm could be changed from brute force to achieve better results, possibly starting at the center of the dictionary and performing a divide and conquer methodology could increase performance overall for any test case.

In addition, the hashing algorithms that were used in this experiment are for the most part outdated. MD5 is not widely used anymore for security because it is considered to be a weak hashing algorithm and easily cracked. Future work could include testing some of the more complex hashing algorithms such as SHA-512 or RIPEMD-320.

# Appendix

## MPI SOURCE CODE

```
/*
 * MPI_Password_Crack.cpp
 *
 * Program designed to utilize the message passing interface (MPI)
 * to parallelize the process of password cracking using the dictionary
 * attack. This will take in a text file of passwords as well as the real
 * password we are looking for, this is to simulate when a hacker has already
 * dumped hashes from the victim machine and is now trying to crack the hash,
 * and perform the appropriate hashing technique to each word in the dictionary.
 * If the hacker is lucky the victim's actual password will be in the dictionary.
 * file chosen and the hacker will be able to crack the hash and compromise the
 * victim's system. Performing this attack serially takes a very long time, so
 * this program facilitates a study to see how much of a speed up a hacker could
 * gain by using MPI.
 *
 * Note:
 * The hashing algorithms used in this study are from zedwood.com
 * Refer to the license file attached.
 *
 * Compile
 * mpic++ md5.cpp sha1.cpp sha256.cpp MPI_Password_Crack.cpp -o
MPI_Password_Crack
 *
 * Run:
 * mpirun -np NUM_PROCESSES MPI_Password_Crack DICTIONARY ACTUAL_PASSWORD
 *
 * Author: Kyle Ray
 * CPE 512 Intro to Parallel Programming
 * Project
 * December 5, 2017
 */

#include <iostream>
#include <vector>
#include <fstream>
#include <cstdlib> // atoi
using namespace std;

// Hashing Function Include
#include "md5.h"
#include "sha1.h"
#include "sha256.h"

// MPI
#include <mpi.h>

// Prototypes
int readDictionary(string filename, std::vector<string>& passwords);
void scatter(std::vector<string>& passwords, std::vector<string>&
local_passwords, int num_passes, int root, int rank, int numtasks);

int main(int argc, char* argv[])
{
    // Variables
    int numtasks, rank, num;
    double tot_finish;
    double start, finish;
```

```

int hash_mode; // 0 - MD5, 1 - SHA-1, 2 - SHA-256
MPI_Status status;

// Start up the MPI Processes
MPI_Init(&argc, &argv); // initialize MPI environment
MPI_Comm_size(MPI_COMM_WORLD, &numtasks); // get total number of MPI
processes
MPI_Comm_rank(MPI_COMM_WORLD, &rank); // get unique task id number

if (rank == 0)
{
    if (argc < 3)
    {
        std::cout << "Usage: mpirun -np numtasks programName hashmode
dictionaryFile passwordToFind" << endl;
        MPI_Abort(MPI_COMM_WORLD, MPI_ERR_INFO);
    }
}

// Store the Hash Mode
hash_mode = atoi(argv[1]);

// Actual Password
std::string actual_pass;
switch (hash_mode)
{
    case 0:
        actual_pass = md5(argv[3]);
        break;
    case 1:
        actual_pass = sha1(argv[3]);
        break;
    case 2:
        actual_pass = sha256(argv[3]);
        break;
}

// Password Container
std::vector<string> passwords;

// Done Flag
bool pass_found = false;

// Total number of passwords, used for sizing
int total_num_passes = 0;

// Root needs to read in and parse the dictionary file
// Should then send the size that each process should allocate to hold their
part
// Scatter the dictionary to other tasks.

if (rank == 0)
{
    std::cout << "Reading Dictionary" << endl;
    if (!readDictionary(argv[2], passwords))
    {
        // Something went wrong with the file
        MPI_Abort(MPI_COMM_WORLD, MPI_ERR_FILE);
    }
}

// Set the total number of passwords to make each process aware
total_num_passes = passwords.size();

```

```

        // Display Size of Dictionary
        std::cout << "Finished reading Dictionary with " << total_num_passes << "
passwords!" << endl;
    }

    MPI_Bcast(&total_num_passes, 1, MPI_INT, 0, MPI_COMM_WORLD);

    // Local Set of Passwords
    std::vector<string> local_passwords;

    // Setup a barrier to make sure that all of the processes enter the region at
the same time
    // for timing analysis
    MPI_Barrier(MPI_COMM_WORLD);

    if (rank == 0)
    {
        std::cout << "Starting the timer. " << endl;
        std::cout << "Scattering...\n";
        start = MPI_Wtime();
    }

    // Scatter the passwords to the processes
    scatter(passwords, local_passwords, total_num_passes , 0, rank, numtasks);

    // Flag for a process to report when done, allows other processes to know
when to stop
    bool someone_done;

    int pass_count = 0;
    double loc_finish = 0;

    // Check each local pass hash and if we find the password then quit and let
all the other processes know
    // Start the main loop
    // 1.) Hash the current password in the MPI process list
    // 2.) Compare the hashed string with the password hash given by the user
    // 3.) If it is a match, then quit and let the other processes know.
    // 4.) Else, process the next password.
    for (int i = 0; i < local_passwords.size(); i++)
    {
        std::string check_pass;
        switch (hash_mode)
        {
            case 0:
                check_pass = md5(local_passwords[i]);    // MD5 algorithm by
zedwood.com
                break;
            case 1:
                check_pass = sha1(local_passwords[i]);    // SHA1 algorithm by
zedwood.com
                break;
            case 2:
                check_pass = sha256(local_passwords[i]); // SHA256 algorithm by
zedwood.com
                break;
        }

        // If we have a match then set the flag
        if (actual_pass == check_pass)
        {

```

```

    pass_found = true;
    loc_finish = MPI_Wtime();
    std::cout << "Pass found by process " << rank << endl
    << "Password is " << local_passwords[i] << endl
    << "Hash is " << check_pass << endl;
}

// Let other processes know if the password is found or not
// TODO: Could probably use non blocking send and receives to do this
somehow, just to make it more efficient.
// Although, all of the loops should be somewhere around the same time
complexity, so it might not matter too much.
// The communication takes too much time because it is executing every
loop.
//MPI_Status status;
//for (int mpitask = 0; mpitask < numtasks; mpitask++)
//{
//    // If not current process, send flag to other processes.
//    if (mpitask != rank)
//    {
//        MPI_Send(&pass_found, 1, MPI_C_BOOL, mpitask, rank, MPI_COMM_WORLD);

//        // Receive flag from other processes
//        MPI_Recv(&someone_done, 1, MPI_C_BOOL, mpitask, mpitask,
MPI_COMM_WORLD, &status);

//        // If password is found by a process, then set exit condition
//        if (someone_done)
//            pass_found = true;
//    }
//}

//pass_count++;

// If the password has been found, no more processing necessary
if (pass_found)
    break;
}

if (!pass_found)
    loc_finish = MPI_Wtime();

MPI_Reduce(&loc_finish, &finish, 1, MPI_DOUBLE, MPI_MIN, 0, MPI_COMM_WORLD);
std::cout << "CLEAN UP: Process " << rank << " is finished processing " <<
endl; //<< pass_count << " number of passwords ending on " <<
local_passwords[pass_count-1] << endl;

MPI_Barrier(MPI_COMM_WORLD);

// Display the results
if (rank == 0)
{
    tot_finish = MPI_Wtime();
    std::cout << "Time to find the password is " << (finish - start) << "
seconds" << endl;
    std::cout << "Total execution time: " << (tot_finish - start) << " seconds"
<< endl;
}

// Terminate MPI Program -- perform necessary MPI housekeeping
// clear out all buffers, remove handlers, etc.

```

```

    MPI_Finalize();
}

// Method to read in the dictionary text file and store the passwords
// into a vector.
int readDictionary(string filename, std::vector<string>& passwords)
{
    // Create file stream object and open the file
    fstream dictFile(filename.c_str());

    // Keep a count of the number of passwords, needed for equal distribution
    int count = 0;

    if (dictFile.fail())
    {
        dictFile.close();
        return count;
    }

    // Read and store each password
    while (!dictFile.eof())
    {
        string line;
        getline(dictFile, line);
        passwords.push_back(line);
        ++count;
    }

    dictFile.close();
    return count;
}

// Method to equally scatter the dictionary to each MPI process
// TODO: Find a faster way to send strings.
// These blocking send and receives take a very long time.
void scatter(std::vector<string>& passwords, std::vector<string>&
local_passwords, int num_passes, int root, int rank, int numtasks)
{
    MPI_Status status;

    // Root handles the distribution of the dictionary file
    if (rank == root)
    {
        int base = num_passes / numtasks;
        int extra = num_passes % numtasks;

        int begin_element = 0;

        for (int mpitask = 0; mpitask < numtasks; mpitask++)
        {
            int num_passes_by_rank = mpitask < extra ? base + 1 : base;
            if (mpitask == root)
            {
                for (int i = 0; i < num_passes_by_rank; i++)
                {
                    local_passwords.push_back(passwords[begin_element]);
                    ++begin_element;
                }
            }
            else
            {
                for (int i = 0; i < num_passes_by_rank; i++)

```



```

        {
            char* pass = const_cast<char*>(passwords[begin_element].c_str());
            MPI_Send(pass, passwords[begin_element].length(),
                     MPI_CHAR, mpitask, mpitask, MPI_COMM_WORLD);
            ++begin_element;
        }
    }
}

// Other processes receive their share of the dictionary
if (rank != root)
{
    int base = num_passes / numtasks;
    int extra = num_passes % numtasks;
    int num_passes_by_rank = rank < extra ? base + 1 : base;
    for (int i = 0; i < num_passes_by_rank; i++)
    {
        // Probe the message
        MPI_Probe(root, rank, MPI_COMM_WORLD, &status);
        int len = 0;
        MPI_Get_count(&status, MPI_CHAR, &len);
        char* buf = new char[len];

        MPI_Recv(buf, len, MPI_CHAR,
                 root, rank, MPI_COMM_WORLD, &status);

        string buf_string(buf, len);
        local_passwords.push_back(buf_string);
    }
}
}

```

## OPENMP SOURCE CODE

```
/*
 * OpenMP_Password_Crack.cpp
 *
 * Program designed to utilize OpenMP parallel library
 * to parallelize the process of password cracking using the dictionary
 * attack. This will take in a text file of passwords as well as the real
 * password we are looking for, this is to simulate when a hacker has already
 * dumped hashes from the victim machine and is now trying to crack the hash,
 * and perform the appropriate hashing technique to each word in the dictionary.
 * If the hacker is lucky the victim's actual password will be in the dictionary
 * file chosen and the hacker will be able to crack the hash and compromise the
 * victim's system. Performing this attack serially takes a very long time, so
 * this program facilitates a study to see how much of a speed up a hacker could
 * gain by using OpenMP.
 *
 * Note:
 * The hashing algorithms used in this study are from zedwood.com
 * Refer to the license file attached.
 *
 * Compile
 * g++ md5.cpp sha1.cpp sha256.cpp OpenMP_Password_Crack.cpp -o
OpenMP_Password_Crack -fopenmp
 *
 * Run:
 * ./OpenMP_Password_Crack HASH_MODE DICTIONARY ACTUAL_PASSWORD NUM_THREADS
 *
 * Author: Kyle Ray
 * CPE 512 Intro to Parallel Programming
 * Project
 * December 5, 2017
 */

#include <iostream>
#include <vector>
#include <fstream>
#include <cstdlib> // atoi
using namespace std;

// Hashing Function Include
#include "md5.h"
#include "sha1.h"
#include "sha256.h"

// OpenMP
#ifdef _OPENMP
#include <omp.h>
#endif

// Prototypes
int readDictionary(string filename, std::vector<string>& passwords);

int main(int argc, char* argv[])
{
    // Variables
    double start, finish;

    if (argc < 5)
    {
        std::cout << "Usage: programName hashMode dictionaryFile passwordToFind
numberOfThreads" << endl;
```

```

    exit(EXIT_FAILURE);
}

int hash_mode = atoi(argv[1]); // 0 - MD5, 1 - SHA-1, 2 - SHA-256

// Actual Password
std::string actual_pass;
switch (hash_mode)
{
case 0:
    actual_pass = md5(argv[3]);
    break;
case 1:
    actual_pass = sha1(argv[3]);
    break;
case 2:
    actual_pass = sha256(argv[3]);
    break;
}

// Password Container
std::vector<string> passwords;

// Done Flag
bool pass_found = false;

// Total number of passwords, used for sizing
int total_num_passes = 0;

// Root needs to read in and parse the dictionary file
// Should then send the size that each process should allocate to hold their
part
// Scatter the dictionary to other tasks.
std::cout << "Reading Dictionary" << endl;
if (!readDictionary(argv[2], passwords))
{
    // Something went wrong with the file
    exit(EXIT_FAILURE);
}

// Set the total number of passwords to make each process aware
total_num_passes = passwords.size();

// Display the number of Passwords
cout << "Finished reading Dictionary with " << total_num_passes << "
passwords!" << endl;

cout << "Starting the Timer " << endl;
cout << "Processing...\n";
// Start the timer
start = omp_get_wtime();

// Check each local pass hash and if we find the password then quit
int num_threads;
if (argc < 5)
    num_threads = 1;
else
    num_threads = atoi(argv[4]);

#pragma omp parallel for num_threads(num_threads) schedule(static, 1)
    for (int i = 0; i < passwords.size(); i++)

```

```

{
    std::string check_pass;
    switch (hash_mode)
    {
    case 0:
        check_pass = md5(passwords[i]);
        break;
    case 1:
        check_pass = sha1(passwords[i]);
        break;
    case 2:
        check_pass = sha256(passwords[i]);
        break;
    }

    // If we have a match then set the flag
    if (actual_pass == check_pass)
    {
        pass_found = true;
        finish = omp_get_wtime();
        std::cout << "Pass found by process " << omp_get_thread_num() << endl
            << "Password is " << passwords[i] << endl
            << "Hash is " << check_pass << endl;
    }

    // Let the other processes know that we are done.
    if (pass_found)
    {
        i = passwords.size();
    }
}

// Display the results
std::cout << "Time to find the password is " << (finish - start) << "
seconds" << endl;

return 0;
}

// Method to read in the dictionary text file and store the passwords
// into a vector.
int readDictionary(string filename, std::vector<string>& passwords)
{
    // Create file stream object and open the file
    ifstream dictFile(filename.c_str());

    // Keep a count of the number of passwords, needed for equal distribution
    int count = 0;

    if (dictFile.fail())
    {
        dictFile.close();
        return count;
    }

    // Read and store each password
    while (!dictFile.eof())
    {
        string line;
        getline(dictFile, line);
        passwords.push_back(line);
        ++count;
    }
}

```

```
}  
  
dictFile.close();  
return count;  
}
```